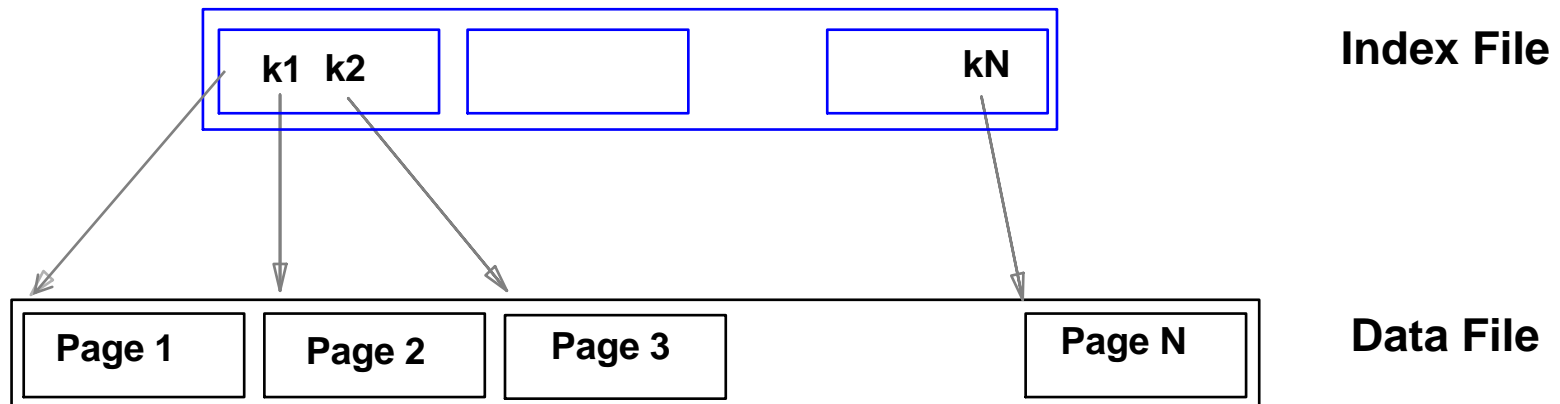


# **Tree-Structured Indexing**

# Range Searches


- ``Find all students with  $3.0 < gpa < 3.5$ ``
  - If records are sorted on gpa, do binary search to find first such student, then scan to find others.
  - Cost of binary search can be quite high.
- Simple idea: Create an `index` file.



☞ *Can do binary search on (smaller) index file!*

# B+ Tree: Most Widely Used Index


- General concept of a tree:



- I. Height of a node: its distance to the root
- II. If a higher level node is connected to a lower level node, then the higher level node is called a parent (grandparent, ancestor, etc.) of the lower level node


# B+ Tree (cont.)

- Balanced tree: all leafs at the same level
  - Example:




- Structure of a B+ tree:
  - It is *balanced*: all leaf nodes at the same level
  - It's node has a special structure

# An example of a B+ tree




# B+ tree: Internal node structure



Each  $P_i$  is a pointer to a child node, each  $K_i$  is a search key value  
Pointers outnumber search key values by *exactly one*.

- **Requirements:**
  - $K_1 < K_2 < \dots < K_n$
  - If the node is not the root, we require  $d \leq n \leq 2d$  where  $d$  is a pre-determined value for this B+ tree, called its *order*
  - If the node is the root, we require  $1 \leq n \leq 2d$
  - For any search key value  $K$  in the subtree pointed by  $P_i$ ,
    - If  $P_i = P_0$ , we require  $K < K_1$
    - If  $P_i = P_1, \dots, P_{n-1}$ , we require  $K_i \leq K < K_{i+1}$
    - If  $P_i = P_n$ , we require  $K_n \leq K$


# B+ tree: leaf node structure



- Each  $r_i$  is a pointer to a record that contains search key value  $K_i$  ...
- L points to the left neighbor, and R points to the right neighbor
- $K_1 < K_2 < \dots < K_n$
- We require  $d \leq n \leq 2d$  where  $d$  is the order of this B+ tree
- We will use  $K_i^*$  for the pair  $K_i, r_i$  and omit L and R for simplicity

# Example: A B+ tree with order of 1

- Each node must hold at least 1 entry, and at most 2 entries




- Given search key values 27, 51, 64, how to find the rids?
  - Search begins at the root, and key comparisons direct it to a leaf



# Example: a B+ tree with order 2


- Search for 5\*, 15\*, all data entries  $\geq 24^*$  ...
- The last one is a range search, we need to do the sequential scan, starting from the first leaf containing a value  $\geq 24$ .



# Cost for searching a value in B+ tree

- In general nodes are pages
- Let  $H$  be the height of the B+ tree: need to read  $H+1$  pages to reach a leaf node
- Let  $F$  be the (average) number of pointers in a node (for internal node, called *fanout* )

- Level 1 = 1 page =  $F^0$  page
- Level 2 =  $F$  pages =  $F^1$  pages
- Level 3 =  $F \times F$  pages =  $F^2$  pages
- Level  $H+1$  = ..... =  $F^H$  pages (i.e., leaf nodes)
- Suppose there are  $D$  data entries. So there are  $D/(F-1)$  leaf nodes
- $D/(F-1) = F^H$ . That is,  $H = \log_F \left( \frac{D}{F-1} \right)$




# B+ Trees in Practice

- Typically, a node is a page
- Typical order: 100. Typical fill-factor: 67%.
  - average fanout = 133 (i.e, # of pointers in internal node)
- Can often hold top levels in buffer pool:
  - Level 1 = 1 page = 8 Kbytes
  - Level 2 = 133 pages = 1 Mbyte
  - Level 3 = 17,689 pages = 133 MBytes
- Suppose there are 1,000,000,000 data entries.
  - $H = \log_{133}(1000000000/132) < 4$
  - The cost is 5 pages read


# Inserting a Data Entry into a B+ Tree

- Find correct leaf  $L$ .
- Put data entry onto  $L$ .
  - If  $L$  has enough space, *done!*
  - Else, must split  $L$  (into  $L$  and a new node  $L2$ )
    - Redistribute entries evenly, put middle key in  $L2$
    - copy up middle key.
    - Insert index entry pointing to  $L2$  into parent of  $L$ .
- This can happen recursively
  - To split an internal node, redistribute entries evenly, but push up middle key. (Contrast with leaf splits.)
- Splits “grow” tree; root split increases height.
  - Tree growth: gets wider or one level taller at top.

# Inserting 16\*, 8\* into Example B+ tree





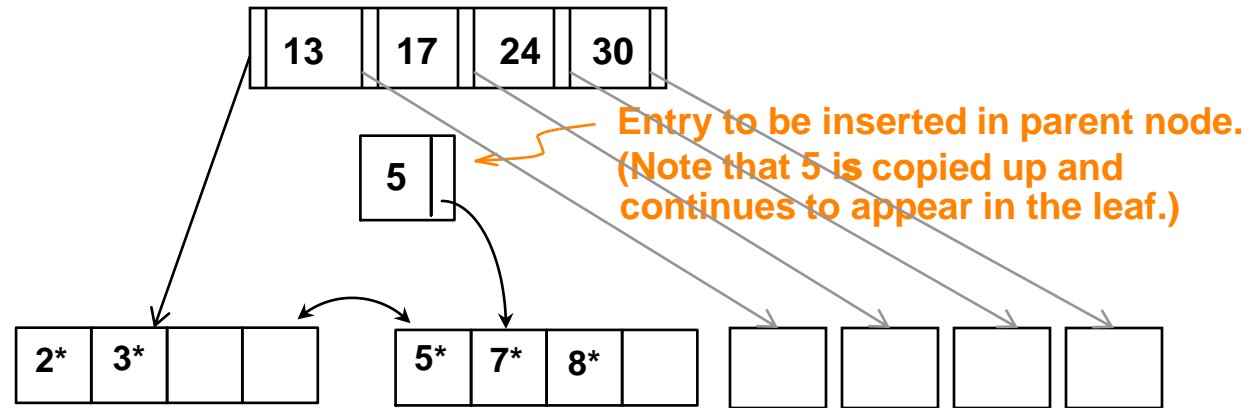
**Overflow!**



One more child generated, must  
add one more pointer to its parent,  
thus one more key value as well.

# Inserting 8\* into Example B+ Tree (order 2)

- Copy the middle value up. Why not push up?
- Observe how minimum occupancy is guaranteed in both leaf and index pg splits.




## Insertion into B+ tree (cont.)




- We delete this pointer!
- But then we should also delete 17
- On the other hand, a value must be inserted into its parent.
- Therefore, we insert 17 to its parent

- This explains why we must push up the middle entry, instead of copying it up, when we split an internal node.



# Example B+ Tree After Inserting 8\*



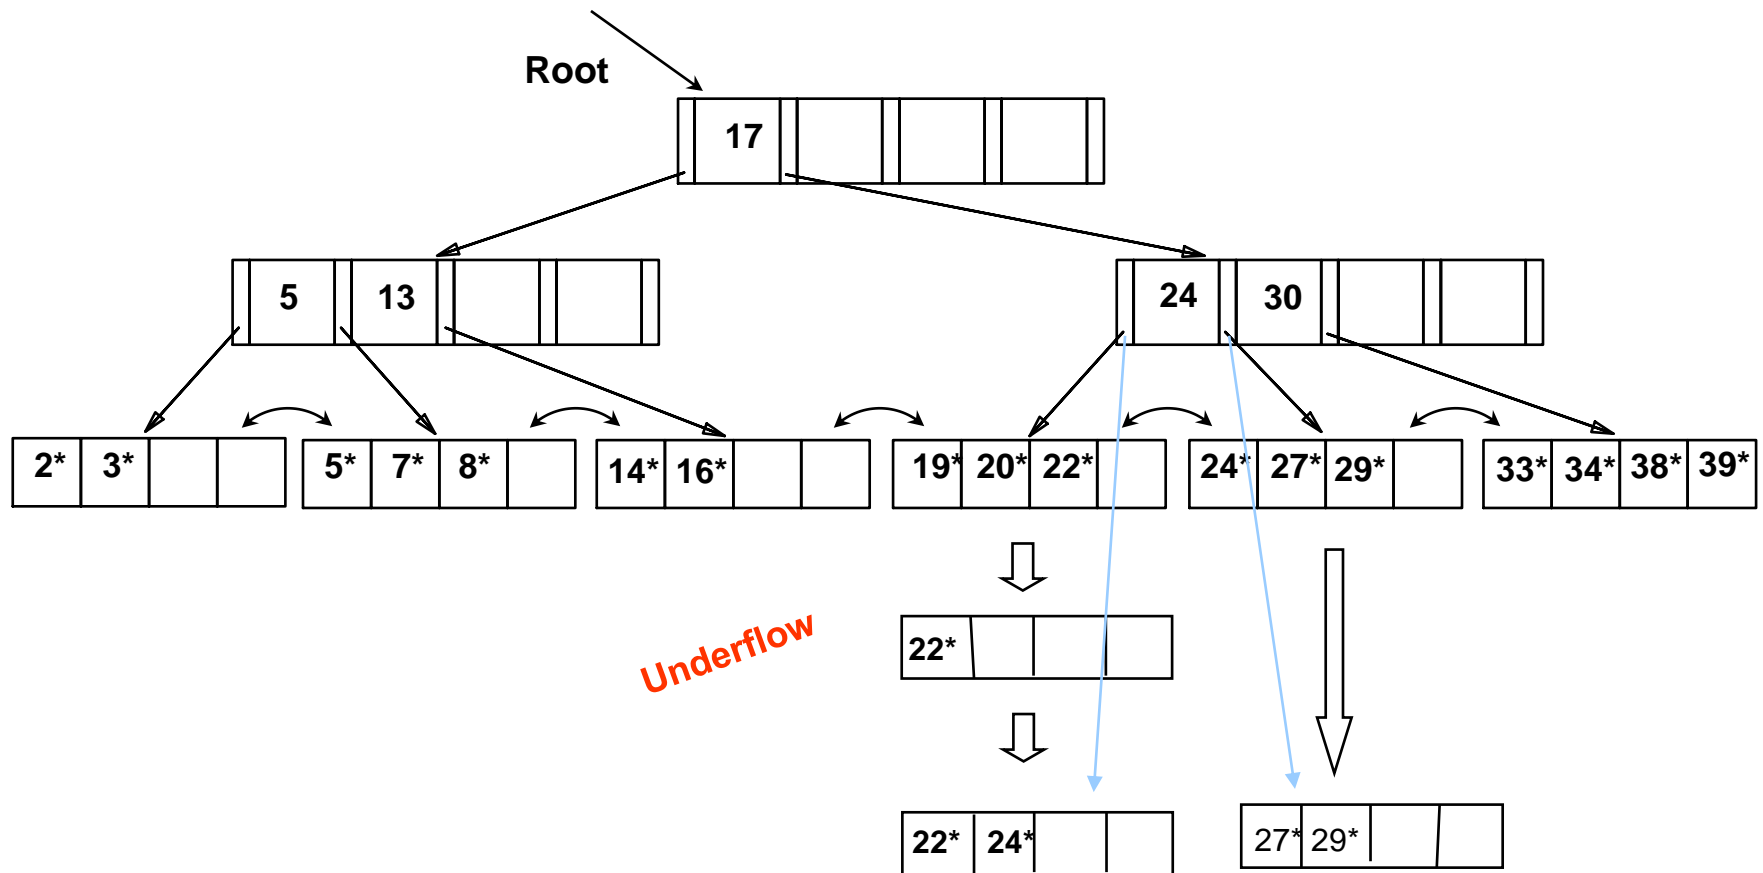
- Notice that root was split, leading to increase in height.
- In this example, we can avoid splitting by re-distributing entries; however, this is usually not done in practice.



# Deleting a Data Entry from a B+ Tree


- Start at root, find leaf  $L$  where entry belongs.
- Remove the entry.
  - If  $L$  is at least half-full, *done!*
  - If  $L$  has only **d-1** entries,
    - Try to **re-distribute**, borrowing from sibling (*adjacent node with same parent as  $L$* ).
    - If re-distribution fails, merge  $L$  and sibling.
- If merge occurred, must delete entry (pointing to  $L$  or sibling) from parent of  $L$ .
- Merge could propagate to root, decreasing height.

# Delete 19\* and 20\*








Have we still forgot something?

# Deleting 19\* and 20\* (cont.)



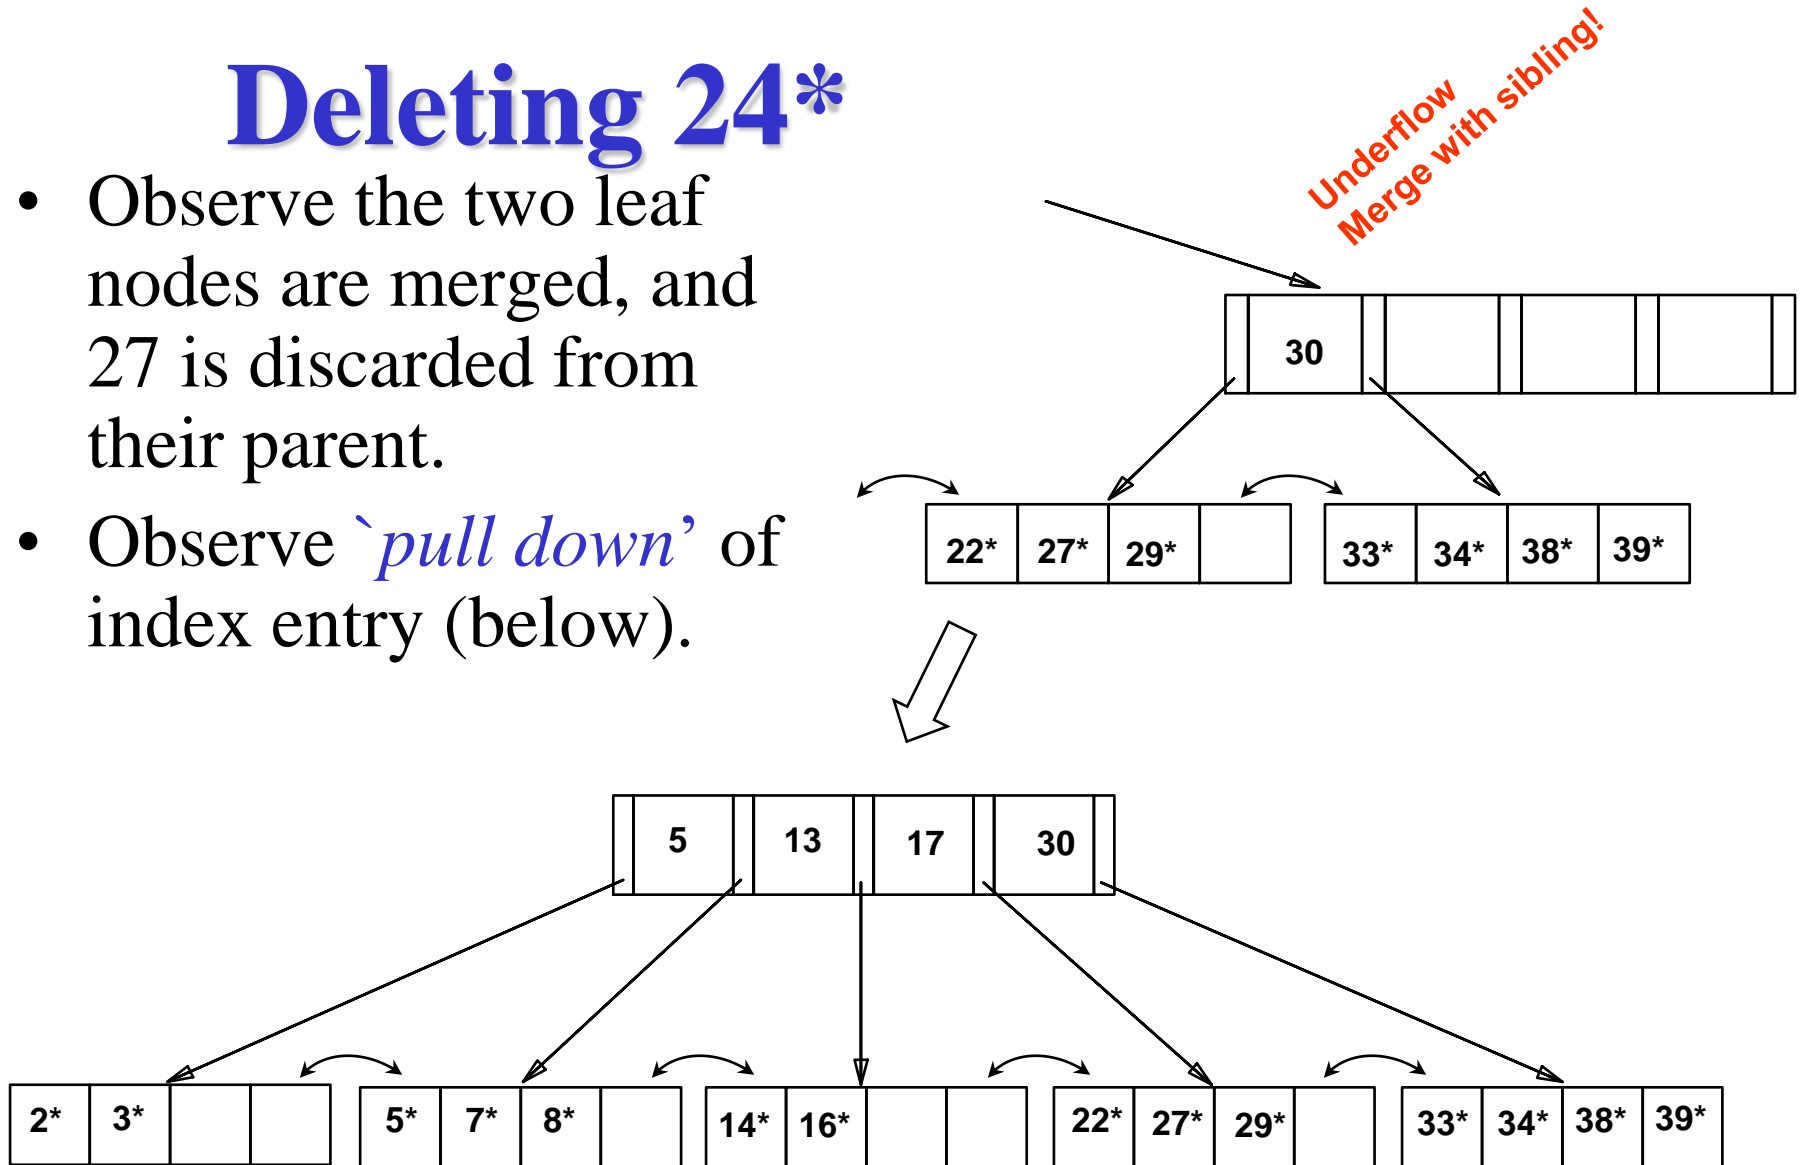
- Notice how 27 is *copied up*.
- But can we move it up?
- Now we want to delete 24
- Underflow again! But can we redistribute this time?






# Deleting 24\*

- Observe the two leaf nodes are merged, and 27 is discarded from their parent.
- Observe '*pull down*' of index entry (below).



# Example of Non-leaf Re-distribution

- Tree is shown below *during deletion* of 24\*. (What could be a possible initial tree?)
- In contrast to previous example, can re-distribute entry from left child of root to right child.



# After Re-distribution

- Intuitively, entries are re-distributed by *'pushing through'* the splitting entry in the parent node.

