


# Hash-based Indexes

# Introduction

- *As for any index, 3 alternatives for data entries  $k^*$* 
  - Data entries are kept in *buckets* (an abstract term)
  - Each bucket is a collection of one primary page and zero or more overflow pages
  - Given a search key value,  $k$ , we can find the bucket where the data entry  $k^*$  is stored as follows:
    - Use a function, called *hash function*, denoted as  $h$
    - The value of  $h(k)$  is the address for the desired bucket (i.e, the address of a bucket is represented by the address of its primary page)
    - $h(k)$  should distribute the search key values uniformly over the collection of buckets
- *Hash-based* indexes are best for *equality selections*. **Cannot** support range searches.
- Static and dynamic hashing techniques exist

# Static Hashing

- # primary pages fixed, allocated sequentially, never de-allocated; overflow pages if needed.
- A simple hash function can be:  $h(k) = f(k) \bmod M$  where  $M = \#$  of buckets and  $f(k) = a \times k + b$
- Example:  $f(k) = k$ . Let  $M = 20$ . Thus  $h(k) = k \bmod 20$ 
  - Assume each page contains two entries
  - For  $k = 1, 21, 41$ , one of them must go to overflow page



# Static Hashing (Cont.)

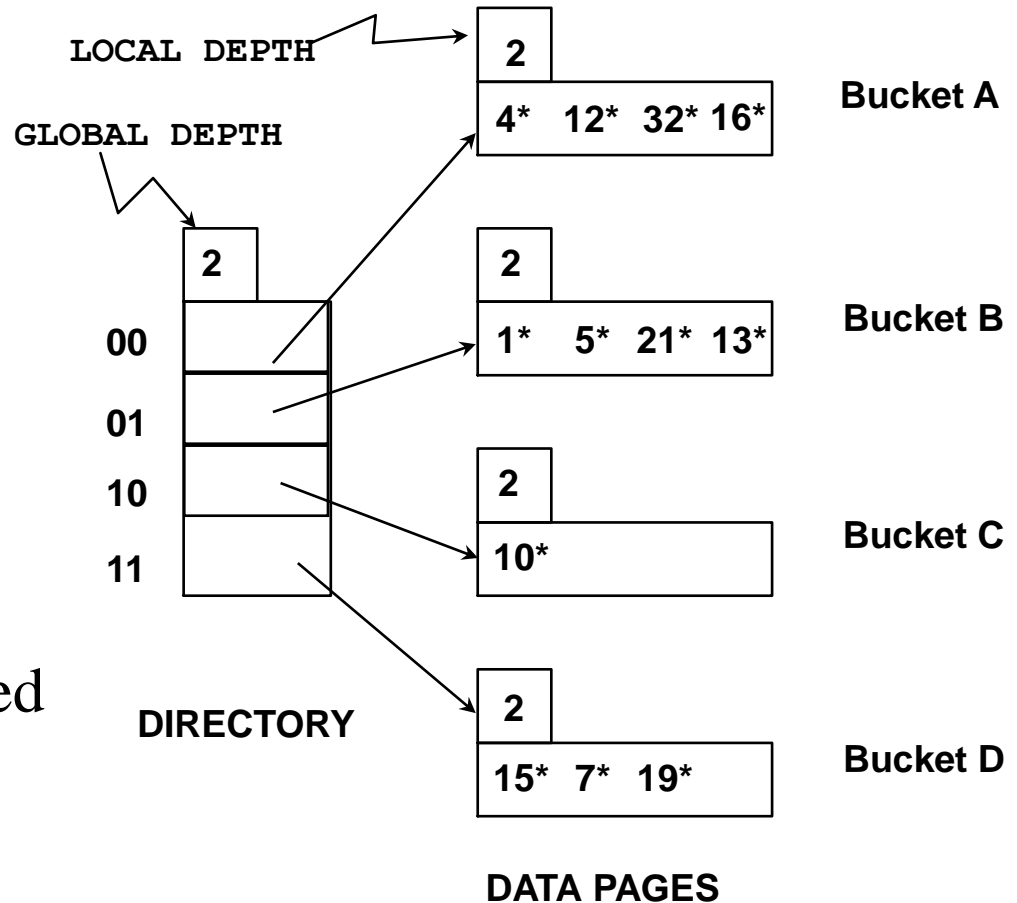
- Buckets contain *data entries*.
- Hash function must distribute values over range of  $0 \dots M-1$ .
- It should be random and uniform.
  - If search key values greatly outnumber  $M$ , then many different key values may be hashed to the same bucket
  - Consider what happens if  $M = 20$  and there are 1000 different search key values:
    - At least one bucket contains 50 values.
    - If the size of a page is 2, then that bucket contains 25 pages: 1 primary and 24 overflow pages
  - Therefore, **long overflow chains** can develop and degrade performance.
  - *Extendible hashing*: Dynamic techniques to fix this problem.

# Extendible Hashing

- Situation: Bucket (primary page) becomes full. Why not re-organize file by *doubling* # of buckets?
  - Must re-hash all data entries to the right buckets
  - Example: assume hash function  $h(k) = k \bmod M$ 
    - For  $M = 4$ , entries  $3^*$  and  $7^*$  both in bucket 3 ( $3 \bmod 4 = 7 \bmod 4 = 3$ )
    - But for  $M = 8$ , entry  $7^*$  will be in bucket 7
  - Can we only re-hash those values that have changed addresses?
    - Difficulties: without re-hashing all the values, we don't know which values keep the old addresses and which get new addresses
  - Reading and writing all pages are expensive!
  - Question: how do we add more buckets, but only re-hash a few data entries?
  - Answer: use a level of indirection, *directory of pointers to buckets*

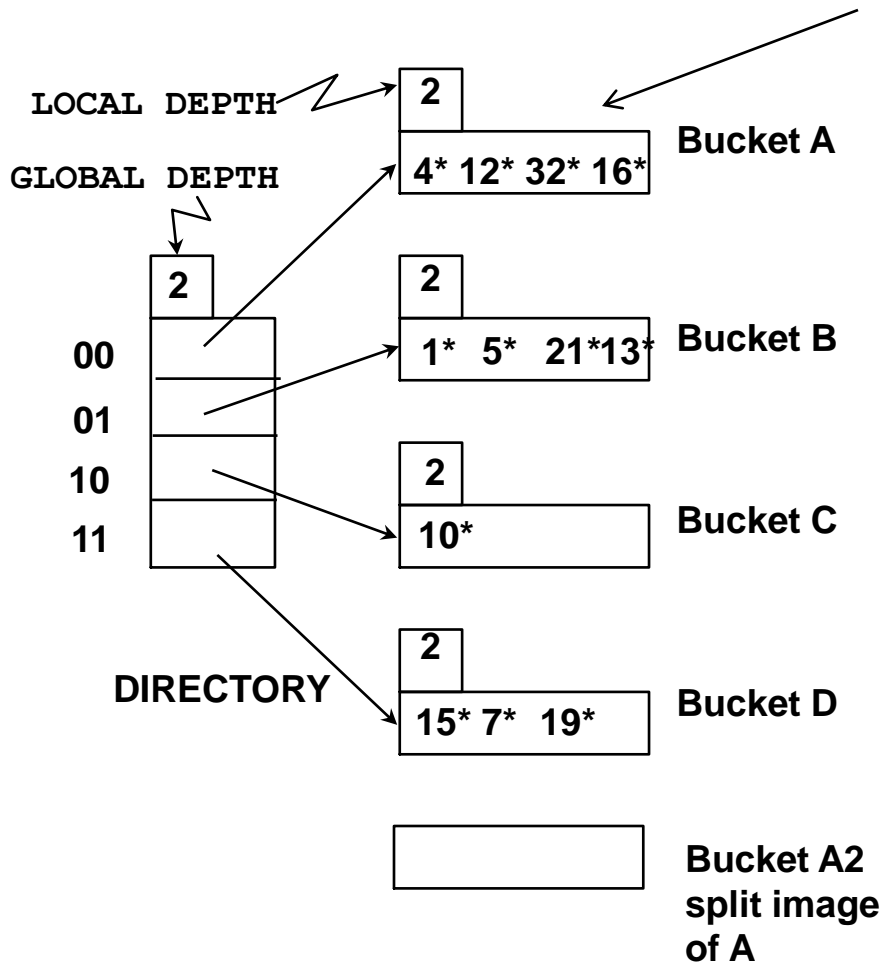
# Example

- $h(r) = r \bmod 32$
- Directory is array of size 4.
- To find bucket for  $r$ , take last '*global depth*' # bits of  $h(r)$ 
  - If  $r = 5$ ,  $h(r) = 5 =$  binary 101, 5\* is in bucket pointed to by 01.



- ❖ Insert: If bucket is full, *split* it (allocate new page, re-distribute).
- ❖ If necessary, double the directory. (As we will see, splitting a bucket does not always require doubling; we can tell by comparing *global depth* with *local depth* for the split bucket.)

# Example (cont.): Insert 20\*



Insert 20\*, causing overflow,  
we do the following:

- Split bucket A to A & A2
- for the five data entries, 4\*, 12\*, 32\*, 16\*, 20\*, if for any  $r^*$  its 3rd bit in  $h(r)$  is 1, then move it to A2:

$$h(4) = 000100$$

$$h(12) = 001100$$


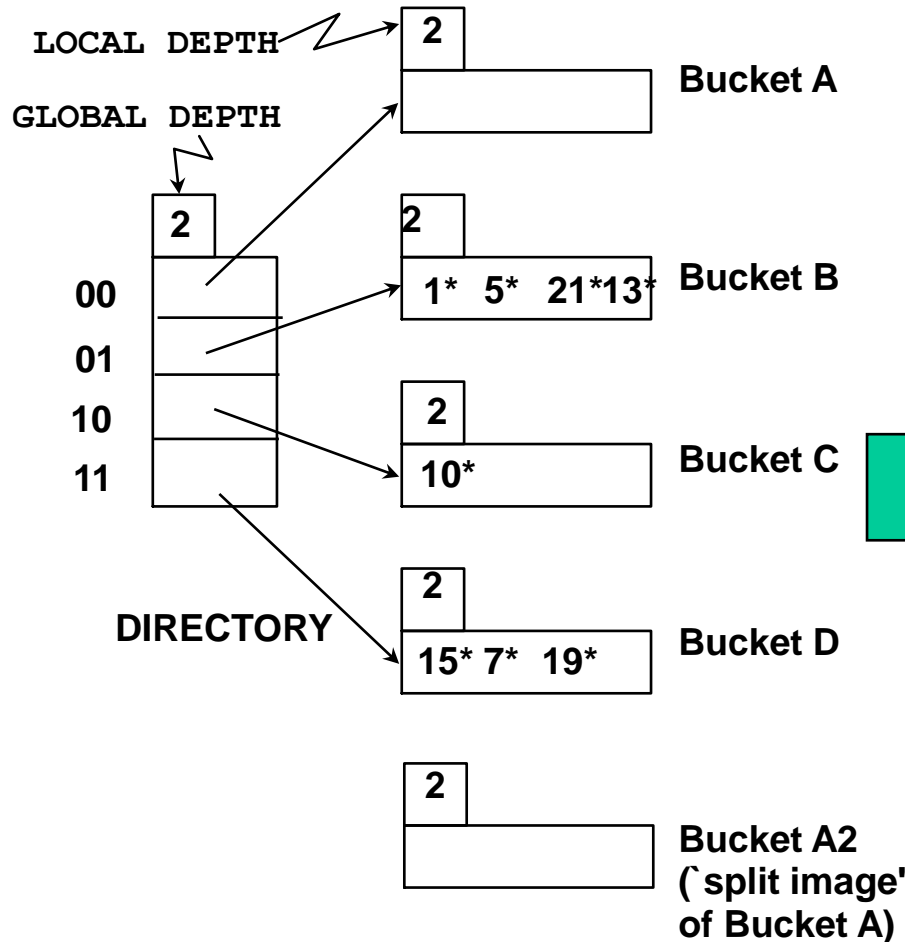
$$h(32) = 100000$$

$$h(16) = 010000$$

$$h(20) = 010100$$




# Insert $h(r)=20$ (Causes Doubling)

$h(16)=010000$   
 $h(32)=000000$   
 $h(20)=010100$   
 $h(12)=001100$   
 $h(4)=000100$





# Insert 9 (Does Not Cause Doubling): $h(9) = 01001$




# Points to Note

- $h(20) = \text{binary } 10100$ . Last **2** bits (00) tell us  $r$  belongs in A or A2. Last **3** bits needed to tell which.
  - *Global depth of directory*: Max # of bits needed to tell which bucket an entry belongs to.
  - *Local depth of a bucket*: # of bits used to determine if the directory need doubling. How?
- When does bucket split cause directory doubling?
  - Before insertion, *local depth* of bucket  $\leq$  *global depth*.
  - After insertion, if overflow, generate split image, and increment the *local depth*
  - If this causes *local depth*  $>$  *global depth*, then directory is doubled, and at the same time increment *global depth*
  - Doubling directory is done by *copying it over* and 'fixing' pointer to split image page. (Use of least significant bits enables efficient doubling via copying of directory!)

# Directory Doubling


## Why use least significant bits in directory?

⇔ Allows for doubling via copying!



## Least Significant

VS.



## Most Significant

# Comments on Extendible Hashing

- If directory fits in memory, equality search answered with one disk access; else two.
  - Directory grows in spurts, and, if the distribution of *hash values* is skewed (e.g., a large number of search key values all are hashed to the same bucket ), directory can grow large.
  - Multiple entries with same hash value cause problems!
- **Delete**: If removal of data entry makes bucket empty, can be merged with 'split image'. If each directory element points to the same bucket as its split image, can halve directory.

# Summary

- Hash-based indexes: best for equality searches, cannot support range searches.
- Static Hashing can lead to long overflow chains.
- Extendible Hashing avoids overflow pages by splitting a full bucket when a new data entry is to be added to it.  
*(Duplicates may require overflow pages.)*
  - Directory to keep track of buckets, doubles periodically.
  - Can get large with skewed data; additional I/O if this does not fit in main memory.