

Recovery

Recoverable

- Transactions may be aborted due to logical failure. e.g. deadlock
- Recoverability is required to ensure that aborting a transaction does not change the semantics of committed transaction's operations.
- Example
 - Write₁(x,2); Read₂(x); Write₂(y,3); Commit₂
 - Not recoverable.
 - T₂ has committed before T₁ commits.
 - The problem is: what can we do if T₁ abort?
 - Delaying the commitment of T₂ can avoid this problem
- A schedule H is called *recoverable* (RC) if, whenever T_i reads from T_j . ($i \neq j$) in H and $c_i \in H$, $c_j < c_i$.
- Intuitively, a history is recoverable if each transaction commits after the commitment of all transactions (other than itself) from which it reads.

Avoiding Cascading Aborts

- Even for recoverable execution, aborting a transaction may trigger further abortions, a phenomenon called *cascading abort*.
- Example
 - Write₁(x,2); Read₂(x); Write₂(y,3); Abort₁
 - T₂ must abort because if it ever committed, the execution would no longer valid.
- A schedule H **avoids cascading aborts** (ACA) if, whenever T_i reads x from T_j ($i \neq j$), $c_j < r_i[x]$.
- That is, a transaction may read only those values that are written by committed transactions or by itself.

Strict Executions


- Avoiding cascading aborts is not always enough from the practical point of view.
- Write₁(x,1); Write₁(y,3); Write₂(y,1); Commit₁; read₂(x); Aborts₂\$
- If we erase the operations from T₂, the resulting execution should be
 - Write₁(x,1); Write₁(y,3); Commit₁
- The value of *y* should be 3.
- The value should be restored for *y* when T₂ is aborted.
- Implement abort by restoring the before images of all Writes of a transaction.
- The **before image** of a *Write(x, val)* operation in an execution is the value of *x* just before this write operation.

- The following example illustrates the problem
 - $\text{Write}_1(x,2); \text{Write}_2(x,3); \text{Abort}_1$
 - Assume the initial value of x is 1.
i.e., the before image of $\text{Write}_1(x,2)$ is 1.
 - Blindly restoring the before image will lead to a wrong result.
- Another example (assume the initial value of x is 1)
 - $\text{Write}_1(x,2); \text{Write}_2(x,3); \text{Abort}_1; \text{Abort}_2$
 - The before image of $\text{Write}_2(x,3)$ is 2.
 - After $\text{Write}_2(x,3)$ has been undone, the value of x should be 1.
- We can avoid these problems by requiring that the execution of a $\text{Write}(x, \text{val})$ be delayed until all transactions that have previously written x are either committed or aborted.
- A schedule H is **strict** (ST) if whenever $w_j[x] < o_i[x]$ ($i \neq j$), either $a_j < o_i[x]$ or $c_j < o_i[x]$ where $o_i[x]$ is $r_i[x]$ or $w_i[x]$.
- This property can be enforced by using the strict two-phase locking protocol (locks are released at the end of transactions), i.e., Strict 2PL guarantees both conflict serializability and strict execution.

- Examples

- $T1 = w1[x] w1[y] w1[z] c1$
- $T2 = r2[u] w2[x] r2[y] w2[y] c2$
- $H1 = w1[x] w1[y] r2[u] w2[x] r2[y] w2[y] c2 w1[z] c1$
- $H2 = w1[x] w1[y] r2[u] w2[x] r2[y] w2[y] w1[z] c1 c2$
- $H3 = w1[x] w1[y] r2[u] w2[x] w1[z] c1 r2[y] w2[y] c2$
- $H4 = w1[x] w1[y] r2[u] w1[z] c1 w2[x] r2[y] w2[y] c2$

- H1 is not RC, because T2 reads y from T1, but $c2 < c1$.
- H2 is RC but not ACA, because T2 has read y from T1 before T1 commits.
- H3 is ACA but not ST, because T2 has overwritten the value written into x by T1 before T1 terminates
- H4 is strict.



Transaction Execution

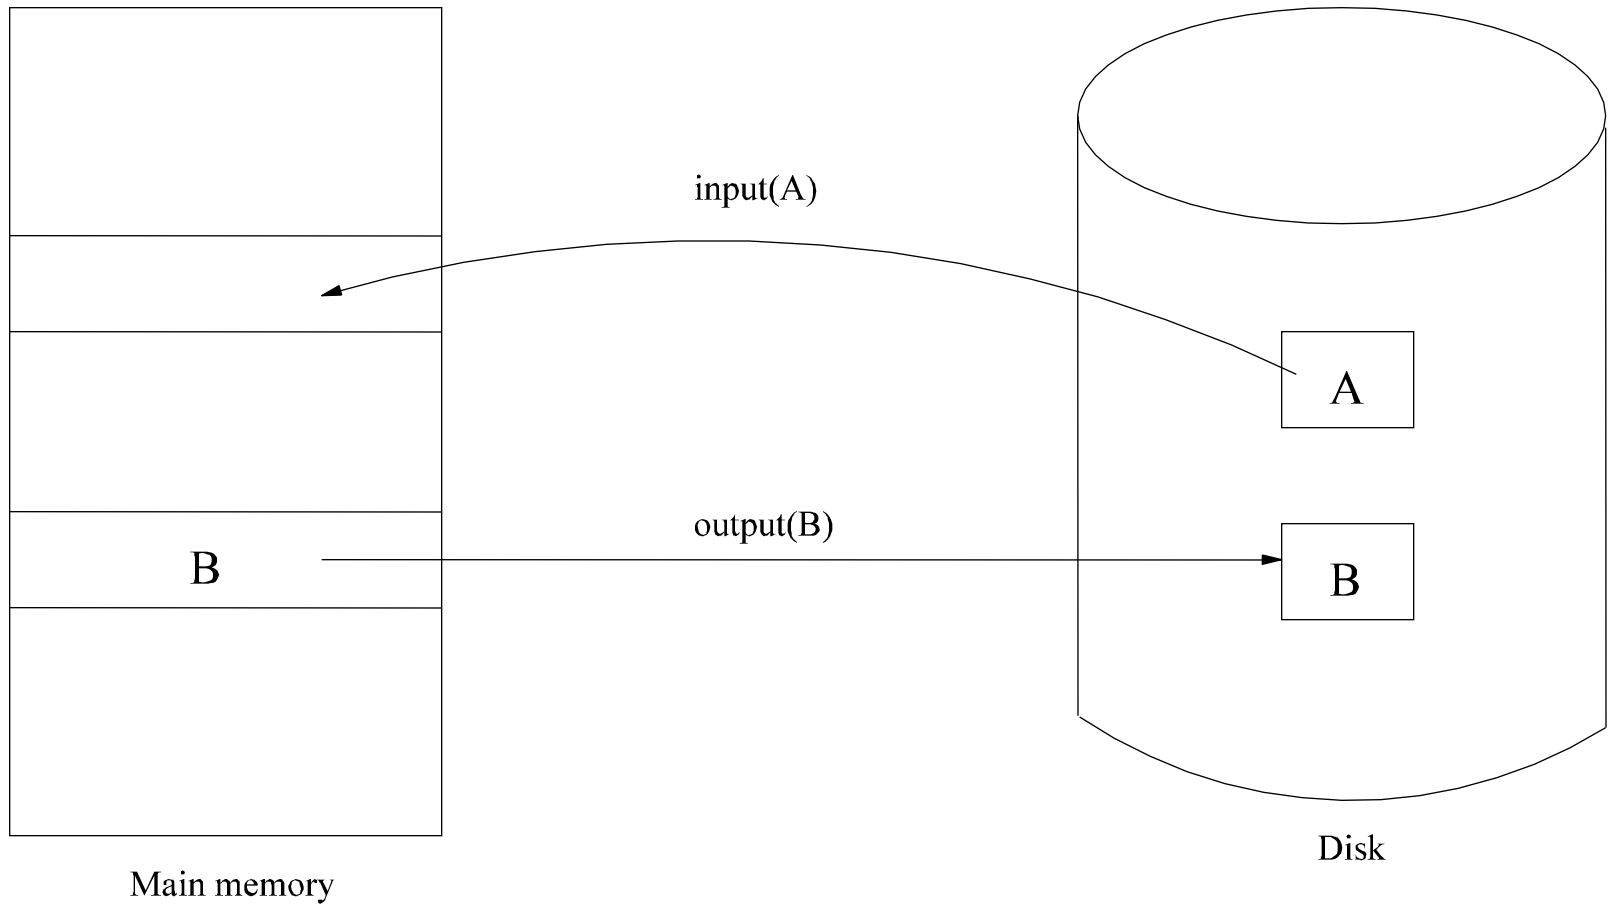
- Assume that **strict two-phase locking protocol** is used.
 - Locks are released at the end of a transaction.
 - If a transaction T_i writes a data item x , no transaction can read or write x before T_i commits or aborts.
 - When a transaction is aborted, the effect of a write operation can be removed by restoring the before image (value of the data item just before the write operation is executed).
 - Recoverable, Avoiding Cascading Aborts, and Strict Execution are guaranteed.

Crash Recovery

- Storage Types
 - Volatile storage: Does not survive system crashes. e.g. main memory, cache memory.
 - Nonvolatile storage: Survives system crashes. e.g. disk, tapes.
 - Stable storage: never lost. e.g. implemented by replication.
- Failure Types
 - Logical errors: cannot continue execution because of internal conditions. e.g. bad input, data not found, overflow, resource limit exceeded.
 - System errors: the system has entered an undesirable state. e.g. deadlock.
 - System crash: hardware malfunctions, causing the loss of the content of volatile storage. e.g. power failure.
 - Disk failure: disk block loses its content. e.g. head crash, failure during a data transfer operation.

- ## Storage Hierarchy

- The database system resides in nonvolatile storage (disk).
- Data transfer are in terms of block (page).
- Physical block: block residing on the disk.
- Buffer block: blocks residing temporarily in main memory.
- Block movements between disk and main memory are initiated through:
 - input(X): From disk to memory
 - output(X): From memory to disk.


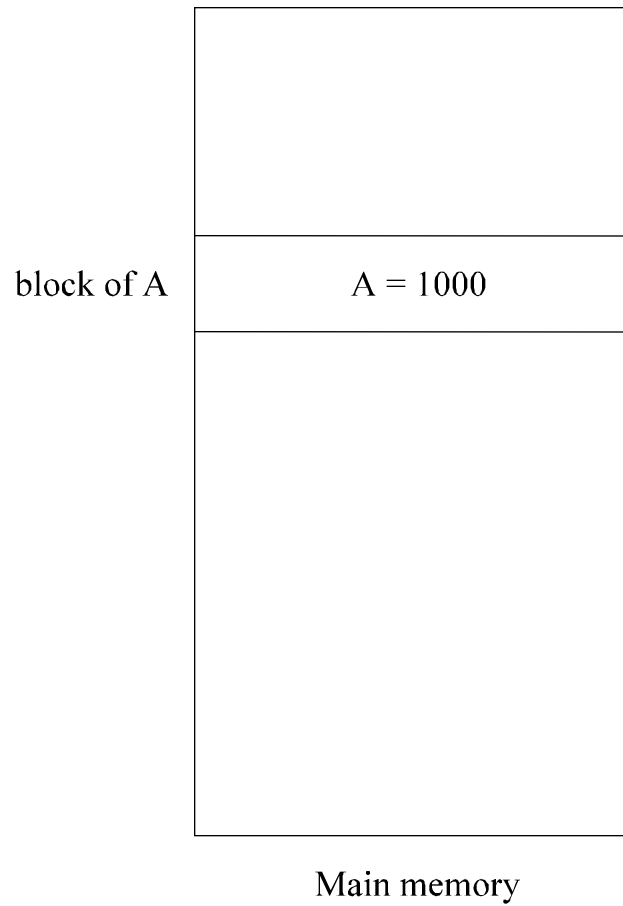


- Transactions interact with the database:
 - Data in database \leftrightarrow program variables.
 - `read(X,xi)`: assigns the value of data item X to the local variable xi.
 - Issue `input(X)`, if X is not in the buffer.
 - Assign the value of X to xi.
 - `write(X,xi)`: assigns the value of local variable xi to data item X.
 - Issue `input(X)`, if X is not in the buffer.
 - Assign the value of xi to X in the buffer block for X.- Not specifically require the transfer of a block from buffer to disk.
- A buffer block is eventually written out to the disk either:
 - Buffer manager needs the memory space.
 - Force-output

Examples


```
read(A,a1)
a1 := a1 - 50
write(A,a1)
read(B,b1)
b1 := b1 + 50
write(B,b1)
```

- The consistency constraint is that the sum of A and B is unchanged.
- Initial values of A and B are \$1000 and \$2000
- Main memory contains the buffer block A.




- `read(A,a1): a1 \leftarrow A.`
- `write(A,a1)`
- `read(B,b1): input(B); b1 \leftarrow B`

```
read(A,a1)
a1 := a1 - 50
write(A,a1)
read(B,b1)
b1 := b1 + 50
write(B,b1)
```




Main memory




Disk

- Write(B,b1)

```
read(A,a1)
a1 := a1 - 50
write(A,a1)
read(B,b1)
b1 := b1 + 50
write(B,b1)
```



Main memory



Disk

What is the state of the database after output(A), but before output(B)?
What if the system crashes at this moment?

Log-Based Recovery

- Recording database modifications in the *log*.
- Each record describes a single database write, and has the following fields:
 - Transaction name.
 - Data item name.
 - Old value (optional)
 - New value
- Special log record:
 - $\langle T_i, \text{start} \rangle$
 - $\langle T_i, \text{commit} \rangle$
 - $\langle T_i, \text{abort} \rangle$

- *Require undo*: If it allows an uncommitted transaction to record in the stable database values it wrote.
- *Require redo*: If it allows a transaction to commit before all the values it wrote have been recorded in the stable database.
- *Undo Rule*: If x 's location in the stable database presently contains the last committed value of x , then that value must be saved in stable storage (*log*) before being overwritten in the stable database by an uncommitted value.
- *Redo Rule*: Before a transaction can commit, the value it wrote for each data item must be in stable storage (*log*).

Immediate Database Modification (undo/redo)

- Allow database modifications to be output to the database while the transaction is active.

T0:

```
Read(A,a1)
a1 := a1 - 50
Write(A,a1)
Read(B,b1)
b1 := b1 + 50
Write(B,b1)
```

T1:

```
Read(C,c1)
c1 := c1 - 100
Write(C,c1)
```

Log:

```
<T0,starts>
<T0,A,1000,950>
<T0,B,2000,2050>
<T0,commits>
<T1,starts>
<T1,C,700,600>
<T1,commits>
```

Note: When T0 and T1 are executed concurrently, their log records may be interleaved.

- Transaction T_i needs to be undone if the log contains $\langle T_i, \text{starts} \rangle$ but does not contain $\langle T_i, \text{commits} \rangle$.
- Transaction T_i needs to be redone if the log contains both $\langle T_i, \text{starts} \rangle$ and $\langle T_i, \text{commits} \rangle$.

Log	Database
$\langle T_0, \text{starts} \rangle$	
$\langle T_0, A, 1000, 950 \rangle$	
	A = 950
$\langle T_0, B, 2000, 2050 \rangle$	
	B = 2050
$\langle T_0, \text{commits} \rangle$	
$\langle T_1, \text{starts} \rangle$	
$\langle T_1, C, 700, 600 \rangle$	
	C = 600
$\langle T_1, \text{commits} \rangle$	

The write operations only update the data items in the buffer.

The updated values may not have been flushed to the hard disk.

<T0,starts>
<T0,A,1000,950>
<T0,B,2000,2050>

Undo T0

<T0,starts>
<T0,A,1000,950>
<T0,B,2000,2050>
<T0,commits>
<T1,starts>
<T1,C,700,600>

Redo T0, undo T1

<T0,starts>
<T0,A,1000,950>
<T0,B,2000,2050>
<T0,commits>
<T1,starts>
<T1,C,700,600>
<T1,commits>

Redo T0, redo T1

Note: Under concurrent execution, the log records from different transactions may interleave.

Detailed Procedure

- *Ti: Start*
 - Write $\langle Ti, start \rangle$ to log
- *Ti: Write(x, v)*
 - If x is not in the buffer, fetch it.
 - Append $\langle Ti, x, ov, v \rangle$ to the log.
 - Write v into the buffer slot occupied by x .
 - Acknowledge the scheduler.
- *Ti: Read(x)*
 - If x is not in the buffer, fetch it.
 - Return the value in x 's buffer slot to the scheduler.

- *Ti: Commit*
 - Write $\langle Ti, commit \rangle$
 - Acknowledge the scheduler.
- *Ti: abort*
 - For each data item x updated by Ti
 - If x is not in the buffer, allocate a slot for it.
 - Copy the before image (ov) of x wrt Ti into x 's buffer slot.
 - Write $\langle Ti, abort \rangle$
 - Acknowledge the scheduler.

- *Restart*

- Discard all buffer slots.
- Let $redone = \{\}$ and $undone = \{\}$.
- Scan the log backward. Repeat the following steps until either $redone \cup undone$ equals the set of all data items, or there are no more log entries. For each log entry $\langle Ti, x, ov, v \rangle$, if $x \notin redone \cup undone$, then
 - if x is not in the buffer, allocate a slot for it.
 - if the commit record of Ti has been found, copy v into x 's buffer slot and set $redone := redone \cup \{x\}$.
 - otherwise, copy the before image (ov) of x wrt Ti into x 's buffer slot and set $undone := undone \cup \{x\}$.
- Acknowledge the completion of Restart to the scheduler.

Example

<T0,starts>
 <T0,A,1000,950>
 <T1, starts>
 <T1,C,700,600>
 <T0,B,2000,2050>
 <T0,commits>
 <T1,A,950,1500>

Suppose after a crash the values of A, B and C found in the stable database are 1500, 2000 and 600 respectively.

← Here is the log found in the stable storage.

From the log, we know that T0 has committed before the crash.

	Action (redo/undo)	A	B	C
		1500	2000	600
<T1,A,950,1500>	undo	950	2000	600
<T0,B,2000,2050>	redo	950	2050	600
<T1,C,700,600>	undo	950	2050	700
<T0,A,1000,950>	no action	950	2050	700

Deferred Database Modification (No-undo/redo)

- Recording all database modifications in the log, but deferring the execution of all write operations until the transaction commits.

Log	Database
<T0,starts>	
<T0,A,950>	
<T0,B,2050>	
<T0,commits>	
	A = 950
	B = 2050
<T1,starts>	
<T1,C,600>	
<T1,commits>	
	C = 600

Before a transaction commits, the update will not be written to the database (not even in the buffer).

After a transaction has committed, the data items in the buffer are updated. The updated values may not have been flushed to the hard disk.

<T0,starts>
<T0,A,950>
<T0,B,2050>

No action is needed

<T0,starts>
<T0,A,950>
<T0,B,2050>
<T0,commits>
<T1,starts>
<T1,C,600>

Redo T0

<T0,starts>
<T0,A,950>
<T0,B,2050>
<T0,commits>
<T1,starts>
<T1,C,600>
<T1,commits>

Redo T0, redo T1

Note: Under concurrent execution, the log records from different transactions may interleave.

Detailed Procedure

- *Ti: Start*
 - Write $\langle Ti, start \rangle$ to log
- *Ti: Write(x, v)*
 - Append $\langle Ti, x, v \rangle$ to the log.
 - Acknowledge the scheduler.
- *Ti: Read(x)*
 - If Ti has previously written into x , then return the after image of x wrt Ti .
 - Otherwise
 - If x is not in the buffer, fetch it.
 - Return the value in x 's buffer slot to the scheduler.

- *Ti: Commit*
 - Write $\langle Ti, commit \rangle$ to the log
 - For each x update by Ti
 - If x is not in the buffer, fetch it.
 - Copy the after image (v) of x wrt Ti into x 's buffer slot.
 - Acknowledge schedule
- *Ti: abort*
 - Write $\langle Ti, abort \rangle$
 - Acknowledge the scheduler.

- *Restart*

- Discard all buffer slots.
- Let $redone = \{\}$.
- Scan the log backward. Repeat the following steps until either $redone$ equals the set of all data items, or there are no more log entries. For each log entry $\langle Ti, x, v \rangle$, if the commit record of Ti has been found and $x \notin redone$, then
 - allocate a slot for x in the buffer;
 - Copy v into x 's buffer slot;
 - $redone := redone \cup \{x\}$.
- Acknowledge the scheduler.

Example

<T0,starts>
 <T0,A,950>
 <T1, starts>
 <T1,C,600>
 <T0,B,2050>
 <T0,commits>
 <T1,A,1500>

Suppose after a crash the values of A, B and C found in the stable database are 950, 2000 and 700 respectively.

← Here is the log found in the stable storage.

From the log, we know that T0 has committed before the crash.

	Action (redo)	A	B	C
		950	2000	700
<T1,A,1500>	No action	950	2000	700
<T0,B,2050>	redo	950	2050	700
<T1,C,600>	No action	950	2050	700
<T0,A,950>	redo	950	2050	700