

CSCI3170 Introduction to Database Systems

Tutorial 9
Concurrency Control and Recovery

Outline

- **Concurrency Control**
 - Transaction and Schedule
 - Conflict Serializability and Precedence Graph
 - View Serializability
 - Shared Lock and Exclusive Lock
- **Recovery**
 - Log-Based Recovery
 - Immediate Database Modification
 - Deferred Database Modification

Serializability and Strict Two-Phase Locking Protocol

CONCURRENCY CONTROL

Transaction

- A sequence of read/write operations
- Examples

T_1 : read(a)
 $a = a + 100$
 write(a)

$T_1 = r_1[a] w_1[a]$

T_2 : $a = 100$
 write(a)
 read(a)

$T_2 = w_2[a] r_2[a]$

Schedule

- A sequence of read/write operations from a set of transactions

- Examples

Suppose there are two transactions, T_1 and T_2

$T_1 = r_1[a] w_1[a]$ $T_2 = w_2[a] r_2[a]$

T1: $r_1[a]$	$w_1[a]$
T2: $w_2[a]$	$r_2[a]$

T1: $r_1[a] w_1[a]$
T2: $w_2[a] r_2[a]$

Possible
schedules



Serial Schedule

- The operations belonging to one single transaction appear together
- Examples

T1: $r_1[a]$ $w_1[a]$

T2: $w_2[a]$ $r_2[a]$

Denoted by T_1T_2

T1: $r_1[a]$ $w_1[a]$

T2: $w_2[a]$ $r_2[a]$



Serializable schedule

- The effect is equivalent to some serial schedules
- Examples

T1: $r_1[a]$	$w_1[a]$
T2: $w_2[b]$	$r_2[a]$



Equivalent to T_1T_2

T1: $r_1[a]$ $w_1[a]$
T2: $w_2[b]$ $r_2[a]$

Conflict Operations

- Two operations are **conflict** if
 - They are operations from different transactions on the same data object
 - At least one of them is a **Write** operation

T1: $r_1[a]$

T2: $w_2[a]$

T1: $w_1[a]$

T2: $r_2[a]$

T1: $w_1[a]$

T2: $w_2[a]$

Conflict Equivalent

- Two schedules S_1 and S_2 are **conflict equivalent** if
 - S_1 and S_2 involve the same operations of the same set of transactions
 - Every pair** of conflicting operations is ordered in the same way in S_1 and S_2

S_1

T1:	$w_1[a]$	$r_1[b]$
T2:	$r_2[b]$	$w_2[a]$

Operations:

$w_1[a], r_1[b], r_2[b], w_2[a]$

Operations on a :

$w_1[a]w_2[a]$ (the same order)

Operations on b :

$r_2[b]r_1[b]$ and $r_1[b]r_2[b]$
(not conflict operations)

S_2

T1:	$w_1[a]$	$r_1[b]$
T2:	$r_2[b]$	$w_2[a]$

Conflict equivalent?



Conflict Serializability

- S_1 is **conflict serializable** if it is conflict equivalent to a serial schedule S_2

S_1

T1:	$w_1[a]$	$r_1[b]$
T2:	$r_2[b]$	$w_2[a]$

Conflict equivalent
(From previous slide)

S_2

T1:	$w_1[a]$	$r_1[b]$
T2:	$r_2[b]$	$w_2[a]$


S_2 is a serial schedule
(Can be denoted by T_1T_2)

Is S_1 Conflict serializable?



Precedence Graph

- Test for conflict serializability
- A directed graph $G=(V,E)$, where
 - V includes all transactions involved in the schedule
 - E consists of all edges $T_i \rightarrow T_j$ for which one of three conditions holds:
 - T_i executes **write(X)** before T_j executes **read(X)**
 - T_i executes **read(X)** before T_j executes **write(X)**
 - T_i executes **write(X)** before T_j executes **write(X)**




Consider the conflict operations

Precedence Graph

- A schedule S_1 is **conflict serializable** iff $G(S_1)$ is **acyclic** (i.e. no cycle)
- If the precedence graph is acyclic, the serialization order (execution order) can be obtained through **topological sorting**

Precedence Graph (3)



- Example

T1:	w₁[a]	w ₁ [b]
T2:	r ₂ [b]	w ₂ [c]
T3:	w ₃ [a]	r ₃ [c]

Operations on **a**:

w₁[a] w₃[a]

Operations on **b**:

r₂[b] w₁[b]

Operations on **c**:

w₂[c] r₃[c]


Conflict serializable?

Yes, it is acyclic. ✓

equivalent to **T₂T₁T₃**

T₂ → T₁ → T₃

(check no edge coming)



View Equivalent

- Two schedules S_1 and S_2 are **view equivalent**
 - If T_i reads the initial value of a data item in S_1 , T_i also reads the initial value of the item in S_2 .
 - If T_i reads an item produced by T_j in S_1 , T_i also reads the item produced by T_j in S_2 .
 - If T_i writes the final value of a data item in S_1 , T_i also writes the final value of the item in S_2 .

T1:	$w_1[a]$	
T2:	$r_2[b]$	$r_2[a]$

T1:	$w_1[a]$	
T2:	$r_2[a]$	$r_2[b]$

In both schedules,

- **T2** reads **b** from initial value
- **T2** reads **a** from **T1**
- The final write of **a** is **T1**

View Serializability

- S_1 is **view serializable** if it is view equivalent to a serial schedule

T1:	$w_1[a]$	
T2:	$r_2[b]$	$r_2[a]$

view serializable?



T2 reads b from initial value

T2 reads a from T1

The final write of a is T1

→ equivalent to $T_1 T_2$

$$\begin{array}{c} T_1 \\ v(a) \end{array} \left| \begin{array}{c} T_1 \\ v(b) \\ v(a) \end{array} \right.$$

T1:	$w_1[a]$	$r_1[b]$
T2:	$w_2[b]$	$r_2[a]$

view serializable?



T1 reads b from T2


T2 reads a from T1

The final write of a is T1

The final write of b is T2

Example (Final Exam 09 Fall Q7)

$w_3[a] \ r_2[a] \ w_1[a] \ w_2[a]$



(a) Please state all the **read-from** relations.

T_2 reads x from T_3

T_2 read a from T_3

(b) Please state which transaction executes the **final write** operation on a .

T_2 issues the final write on x

(c) Is the history **view serializable**? Why?


Please state the **serialization order** if it is view serializable.

Yes. The history has the same read from relation and same final write operation as $T1 \ T3 \ T2$

Example (Final Exam 09 Fall Q7)

$w_3[a] r_2[a] w_1[a] w_2[a]$

(d) Draw the serialization graph of the above history.



(e) Is the history conflict serializable? Why?

Please state the serialization order if it is conflict serializable.

No, because there is a cycle in the serialization graph.

Strict Two-Phase Locking Protocol

- If a transaction wants to **read/modify** an object, it first requests a **shared/exclusive** lock on the object.
- All locks held by a transaction are released when the transaction is completed.
- **Pro:** Allow only conflict serializable schedules.
- **Con:** Deadlock may occur.

Shared Lock and Exclusive Lock

- Shared Lock - **lock-S**
 - Used when the transaction only reads the data object
 - **Allow read** (for other transactions)
 - **Not allow write** (for other transactions)
- Exclusive Lock - **lock-X**
 - Used when the transaction has a write operation on the data object
 - **Not allow read** (for other transactions)
 - **Not allow write** (for other transactions)

T1: $r_1[a]$ $r_1[b]$ $w_1[b]$



T1: **lock-S(a)** $r_1[a]$ **lock-X(b)** $r_1[b]$ $w_1[b]$ **unlock(a)** **unlock(b)**

Deadlock

T1	T2
lock-S(a)	
	lock-S(b)
	r ₂ [b]
r ₁ [a]	
lock-X(b)	
	lock-X(a)
	r ₂ [a]
	w ₂ [a]
r ₁ [b]	
w ₁ [b]	
unlock(a)	
unlock(b)	
	unlock(b)
	unlock(a)

T₁ waits for the lock of b (i.e. T₁ waits for T₂ to release the lock of b)

T₂ waits for the lock of a (i.e. T₂ waits for T₁ to release the lock of a)

Log-Based Recovery for Immediate and Deferred Database Modifications

RECOVERY

Log-Based Recovery

- Recording database modifications in the **log**
- Log record
 - **<transaction_name, start>**
 - **<transaction_name, item_name, new_value>**
 - **<transaction_name, commit>**

T0:

Read(A,a1)
a1 := a1 - 50
Write(A,a1)
Read(B,b1)
b1 := b1 + 50
Write(B,b1)

Log:

<T0,starts>
<T0,A,950>
<T0,B,2050>
<T0,commits>

Guarantees

Values have to be **updated** if the transaction is **committed**

Values have to be **restored** if the transaction is **uncommitted**

Immediate Database Modification

- Allow database modifications to be output to the database while the transaction is active

Log record	Database process
<T0,starts>	A = 300, B = 1000
<T0,A,950>	A = 950, B = 1000
<T0,B,2050>	A = 950, B = 2050
<T0,commits>	A = 950, B = 2050

It is possible that the updated values have not been flushed to the hard disk yet.

Recovery

Redo the transaction if the transaction is **committed**

Undo the transaction if the transaction is **uncommitted**

Deferred Database Modification

- Recording all database modifications in the log, but deferring the execution of all write operations until the transaction commits

Log record	Database process
<T0,starts>	A = 300, B = 1000
<T0,A,950>	A = 300, B = 1000
<T0,B,2050>	A = 300, B = 1000
<T0,commits>	A = 950, B = 2050

It is possible that the updated values have not been flushed to the hard disk yet.

Recovery

Redo the transaction if the transaction is **committed**